

DATA CARD

Input vs. Output	• What are users inputting (for example, audio, text, images)? • What are the outputs or labels your model will predict or generate?	Your response: Input: Text Output: one chosen (out of the ten) predefined category
Classification problem	• What is the classification task, if any? • What categories or classes are you working with, and why?	Your response: Classification of complaints into 10 predefined categories: billing, product_defect, delivery_shipping, refund_return, customer_service, account_access, fraud_unauthorized, warranty_repair, subscription_cancel, general_inquiry. These categories were provided by the competition's organisers.
Labeling	• Where do your labels come from—domain experts, community members, crowd workers, AI-generated sources? • How will you check for consistency and accuracy?	Your response: The labels come from the Competition's organisers. Our team will check label consistency indirectly through manual spot-checks of a sample of training examples against their assigned category.

Data creation	• How will you collect your data? • Will it come from direct sources (like interviews or recordings), or will you synthesize data in low-resource settings (e.g., for under-resourced languages)?	Your response: Primary training data is provided directly by the competition (380 labelled examples, 160 unlabeled test examples). The team is still deciding whether to supplement this with augmented versions of existing labelled examples (e.g., paraphrasing or back-translation) to address class imbalance and limited volume.
Web collection ethics	• Do you plan to source data from the web? • If yes, which sites or platforms? • What tools will you use - AI tools, human workers, or both? • Have you addressed issues of consent, copyright, and privacy?	Your response: To be decided, if needed
Bias and representation	• Who might be left out of your dataset? • How will you ensure that the dataset reflects the diversity of the community and includes underrepresented voices?	Your response: The provided dataset's demographic and geographic origin is not disclosed, so it's unclear whose complaints it reflects. Given the dataset's small size (380 examples across 10 categories), rarer categories are likely

		underrepresented, risking weaker model performance precisely where errors may carry more real-world consequence (e.g., fraud_unauthorized, account_access).
Values and community alignment	• How does your dataset design reflect the values you identified in 01 Build Your Own Small Language Model • How will your community benefit from this dataset, and how will you ensure accountability to them?	Your response: The dataset design should reflect the values guiding this project: accuracy with fairness across categories, timeliness in resolving consumer concerns, and trust in preserving human judgment for sensitive or ambiguous cases. Because the training data is small and its provenance is not fully known, the team commits to transparency about these limitations rather than overstating the dataset's representativeness or the model's real-world readiness.