

Stakeholder Groups

Consumers

Who they are: Individuals who submit support complaints into the system, spanning the full range of issue types the model classifies, from routine billing questions to urgent fraud or account access reports.

Why they are important to the project: They are the direct end users of the classification outcome; the entire purpose of the project, faster and more consistent routing, is measured by whether it actually improves their experience. A model that is technically accurate but fails consumers in practice (e.g., mishandling urgent cases) would undermine the project's core goal.

What matters most to them: Being routed correctly and quickly, especially when their complaint is urgent, since delay can mean ongoing financial or account exposure. They also want fair treatment regardless of how they phrase their complaint, and reassurance that a human is still accountable for sensitive decisions.

Support Staff / Frontline Agents

Who they are: The employees who currently manually sort and resolve complaints, and who will use the model's classification output (or work alongside it) once deployed.

Why they are important to the project: They are the operational users who determine whether the tool succeeds in practice. Their day-to-day exposure to real complaints gives them the clearest view of where the model succeeds or fails, and their willingness to trust and adopt the tool directly affects whether it delivers its intended efficiency gains.

What matters most to them: Reducing manual sorting workload without undermining their expertise or job security. They want classifications they can trust enough to act on, an easy way to override mistakes, and confidence that the system supports their judgment rather than replacing it.

Potential Benefits and Risks

Consumers

Benefits: Faster resolution, not having to repeat their issue across departments, and higher-stakes complaints reaching the right team sooner.

Harms/concerns: Misclassification of urgent complaints (e.g., fraud routed as "general_inquiry") causing real delay and harm; unequal accuracy for people who phrase complaints in less "typical" ways; loss of trust if errors go unexplained or unappealable.

Ethical, social, cultural, and economic considerations:

- Ethical: fairness in error distribution — errors should not fall disproportionately on any group of consumers based on how they write or what language they use.
- Social: risk of eroding trust in the support system if consumers feel dismissed or misrouted without recourse.
- Cultural: the dataset's undisclosed demographic and geographic origin makes it unclear whether diverse phrasing styles, dialects, or cultural expressions of complaints are represented.
- Economic: for consumers, delayed handling of financial or fraud-related complaints carries direct monetary risk.

Support Staff

Benefits: Less time spent on manual sorting, more time on actual resolution, and a tool that surfaces edge cases for review rather than hiding them.

Harms/concerns: Over-reliance on model output leading to complacency or missed errors; erosion of triage expertise over time; anxiety that the tool signals eventual job displacement rather than support.

Ethical, social, cultural, and economic considerations:

- Ethical: responsibility for high-stakes misclassifications should remain shared between the model and accountable staff, not silently absorbed by the system.
- Social: staff morale and sense of professional value could be affected if the tool is introduced without their input.
- Economic: workload reduction could be framed positively (more time for resolution) or negatively (justification for staffing cuts), depending on how the tool is positioned and deployed.

Engagement Method

Consumers — Distributed Dialogue

Why the method is appropriate: Consumers are a large, dispersed population with brief, one-off contact with the system at the moment they submit a complaint. A method built for lightweight, distributed, asynchronous input across many people fits this reality far better than a workshop or committee model.

Why it is suitable for this group: It respects that most consumers want their issue resolved, not a seat on a panel, while still gathering signal on whether people whose complaints don't fit "typical" phrasing are being underserved.

How the engagement would be conducted: A short, optional feedback prompt would appear after a complaint is routed and resolved, asking whether the consumer felt it reached the right team and was handled with appropriate urgency, alongside a clear, low-friction path to flag a suspected misclassification. Responses would be aggregated periodically (e.g.,

monthly) rather than reviewed case-by-case, to identify patterns across many consumers rather than relying on individual complaints.

Support Staff — User Committee

Why the method is appropriate: Support staff are a bounded, identifiable group embedded in the organisation who can be convened repeatedly over the life of the project, unlike consumers who interact with the system only briefly.

Why it is suitable for this group: Staff have sustained, hands-on exposure to the model's outputs every working day, making them uniquely positioned to spot recurring failure patterns. A committee gives them real, ongoing influence over how the tool evolves rather than having it imposed on them.

How the engagement would be conducted: A small standing committee of frontline agents (rotating membership to avoid burnout) would meet regularly (e.g., biweekly or monthly) to review a sample of recent classifications, flag systematic misclassifications, and discuss any proposed changes to categories or thresholds before they are deployed. This formalises the manual spot-checking already planned for label verification into a structured, accountable process with real decision input.

Influence on Project Design

Dataset: Consumer distributed dialogue would surface complaint phrasing styles or demographics underrepresented in the current 380-example dataset, informing targeted data collection or augmentation. Support staff feedback would identify systematically confused category pairs, guiding where additional labelled examples are most needed.

Model design: Recurring committee review could inform decisions such as adjusting classification thresholds for high-stakes categories (e.g., requiring higher confidence before auto-routing fraud_unauthorized) or flagging low-confidence predictions for mandatory human review rather than automatic routing.

Outputs: Feedback could shape how classification results are presented, for example, showing staff a confidence score or top-two possible categories for ambiguous cases, rather than a single forced label, so staff judgment remains meaningfully part of the decision.

Evaluation: Both groups' feedback would push evaluation beyond aggregate accuracy toward per-category and per-group performance metrics, ensuring the model is not judged as "successful" while systematically underperforming on rarer or higher-stakes categories.

Deployment decisions: Concerns from both groups support a staged deployment: initially running the model in an assistive capacity (suggesting a category for staff to confirm) rather than fully automated routing, especially for sensitive categories, with automation expanded only as trust and evidence of reliability accumulate.